

Retail Sales Performance Tracker

Technical Documentation

1. Project Overview

This document provides an in-depth technical explanation of the Retail Sales Performance Tracker, which is an end-to-end analytics pipeline that ingests retail data, transforms it for analysis, and visualizes KPIs through an interactive dashboard. It includes engineering decisions, setup steps, configuration files, code structure, testing logic, CI/CD workflow, and encountered issues. The solution simulates how modern data teams build business intelligence systems using modular, cloud-friendly, and open-source tools.

2. Project Goals

- Ingest real retail sales data from a CSV file
- Load the cleaned data into a cloud-based warehouse (Snowflake)
- Build staging and fact models using dbt
- Aggregate sales, profit, and order metrics by time and category
- Apply time-series transformations (rolling average, YTD)
- Visualize business metrics in a Streamlit dashboard
- Add CI/CD validation using GitHub Actions

3. Dataset Summary

- Source: Kaggle – [Superstore Sales Dataset](#)
- Structure: Tabular CSV with ~10,000 rows
- Fields include:
 - Order ID, Customer ID, Order Date, Ship Date
 - Product Name, Category, Sub-Category
 - Region, State, Country, City, Postal Code
 - Sales, Profit, Quantity, Discount

This data is highly relatable and relevant for demonstrating retail business analytics.

4. Technology Stack and Why Each Was Chosen

Component	Tool/Service	Justification
Data Warehouse	Snowflake	Scalable, supports SQL-based modeling, free trial credits
Data Transformation	dbt (data build tool)	Enables modular, tested SQL pipelines with documentation
Scripting	Python	Industry standard for automation and ingestion

Dashboarding	Streamlit	Easy-to-build, interactive, Python-native dashboards
Version Control	Git + GitHub	Standard practice for team collaboration
CI/CD Automation	GitHub Actions	Cloud-native pipeline automation without cost

5. File and Folder Structure

Retail-Sales-Performance-Tracker/

```

├── data/                # Raw CSV file (excluded from Git)
├── dbt_project/         # dbt config, models, schema.yml
├── diagrams/           # Dashboard preview screenshots
├── scripts/            # Data upload Python script
├── streamlit_dashboard/ # Dashboard + Snowflake connector
├── .github/workflows/   # CI pipeline YAML for dbt run/test
├── .env                # Snowflake credentials (ignored)
├── .gitignore
├── README.md
└── requirements.txt

```

6. Data Ingestion: Python to Snowflake

File: `scripts/upload_to_snowflake.py`

Workflow:

1. Read the CSV into a pandas DataFrame
2. Standardize column names (e.g., Order ID → order_id)
3. Connect to Snowflake using `snowflake.connector`
4. Write the DataFrame into RAW.RAW_SUPERSTORE using `write_pandas()`

Code Highlights:

```

from snowflake.connector.pandas_tools import write_pandas
...
df.columns = df.columns.str.strip().str.lower().str.replace(" ", "_")
write_pandas(conn, df, table_name='RAW_SUPERSTORE')

```

.env Example:

```

SNOWFLAKE_USER=your_username
SNOWFLAKE_PASSWORD=your_password
SNOWFLAKE_ACCOUNT=goyjbep-ib68545
SNOWFLAKE_WAREHOUSE=DEMO_WH
SNOWFLAKE_DATABASE=SUPERSTORE_DB
SNOWFLAKE_SCHEMA=RAW

```

Issues & Resolutions

- `UnicodeDecodeError - Used encoding='ISO-8859-1'`
- `invalid identifier` - Removed quotes and renamed columns cleanly
- `pyarrow version mismatch` - Installed compatible version or downgraded

7. dbt Modeling (Transformations)

Models Built

File	Description
<code>stg_orders.sql</code>	Clean staging of raw orders
<code>fct_sales.sql</code>	Aggregated metrics by category and region
<code>fct_sales_by_month.sql</code>	Monthly breakdown
<code>fct_sales_rolling_avg.sql</code>	3-month rolling avg using window functions
<code>fct_sales_by_month_ytd.sql</code>	YTD sales and profit cumulative model

Testing via `schema.yml`

```
- name: fct_sales
  columns:
    - name: category
      tests:
        - not_null
    - name: total_sales
      tests:
        - not_null
```

Documentation & Lineage

```
dbt docs generate
dbt docs serve --port 9000
```

- Visualizes table dependencies
- Describes models and test coverage

Issues & Resolutions

- `dbt not detecting profile` - Recreated `~/.dbt/profiles.yml`
- `dbt Cloud blocking CLI` - Switched to `dbt-snowflake` via pip
- `Port 8080 busy` - Used `--port 9000`

8. Streamlit Dashboard

Main File: `streamlit_dashboard/app.py`

Features Implemented:

- **KPI Cards:** Total Sales, Total Profit, Total Orders (via `st.metric()`)
- **Sales by Category:** Bar chart using `groupby()`
- **Rolling Average:** 3-month average via `fct_sales_rolling_avg`
- **YTD Sales:** Cumulative line chart from `fct_sales_by_month_ytd`

Helper File: `snowflake_connector.py`

- Uses `load_dotenv()` to read `.env`
- Secure Snowflake connection
- Queries raw/fact/rolling/ytd tables separately

Issues & Fixes

- `NoneType` has no attribute `find` - Fixed `.env` not loading
- Chart missing - Ensured Streamlit reloads after changes

9. GitHub Actions: CI/CD Workflow

File: `.github/workflows/dbt_ci.yml`

Trigger: Manual (`workflow_dispatch`)

Workflow Steps:

- Setup Python 3.10
- Install `dbt-snowflake`
- Load `DBT_PROFILE` from GitHub Secrets
- Run `dbt debug`, `dbt run`, `dbt test`

GitHub Secret Used: `DBT_PROFILE` → full contents of `profiles.yml`

Issues & Fixes


- `push rejected` - Pulled remote first: `git pull origin main --allow-unrelated-histories`
- Profile secret not found - Re-pasted clean YAML into GitHub secrets

10. Dashboard Preview

Visuals Displayed:

- KPI cards
- Bar chart: Sales by Category

- Line chart: 3-Month Rolling Avg
- Line chart: YTD Sales



11. Testing Strategy

dbt Tests

- Defined in `schema.yml`
- Test types: `not_null`, `unique`
- Covers primary keys, aggregates, and time columns

Manual Checks

- Compared outputs between models
- Validated counts in Streamlit vs Snowflake

Repository: [Retail-Sales-Performance-Tracker](#)